

Contents

- 1 Executive Summary
- 2 Introduction
- 3 Importing Libraries
- 4 Loading Data
- 5 EDA
- 6 Construct a Linear Regression Model
  - 6.1 Residual Plot
  - 6.2 Check in the Independence
- 7 Modelling (treating) County as a Fixed Effect
  - 7.1 Checking the Coefficients
- 8 Build a Mixed Effect Model
- 9 The Nested Random Effects
- 10 Conclusion

1 Executive Summary

Mixed effect models are a type of regression model that take into account both (1) variation that is explained by the independent variables of interest (like age) – fixed effects, and (2) variation that is not explained by the independent variables of interest (like county)– random effects. Since the model includes a mixture of fixed and random effects, it's called a mixed model.

2 Introduction

Capturing the average time spent on a website

Improving a website's bounce rate and average time on a page will ultimately improve SEO (search engine optimization), The goal in this project is to find out if "Age" of a website visitor has any relationship with the average time spent on the website. For this purpose we have decided to perform a linear regression model and explain its findings.

3 Importing Libraries

```
In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import sklearn as sk
from math import sqrt
import warnings
warnings.filterwarnings('ignore')
```

4 Loading Data

```
In [3]: data=pd.read_csv("Downloads/data.csv")
```

```
In [4]: data.head(5)
```


Out[4]:

	bounce_time	age	county	location
0	165.548520	16	devon	a
1	167.559314	34	devon	a
2	165.882952	6	devon	a
3	167.685525	19	devon	a
4	169.959681	34	devon	a

5 EDA


```
In [10]: sns.distplot(data.bounce_time)
```

Out[10]: <matplotlib.axes.\_subplots.AxesSubplot at 0x1416a1cee08>



```
In [13]: sns.distplot(data.age)
```

Out[13]: <matplotlib.axes.\_subplots.AxesSubplot at 0x1416a83e908>



```
In [15]: from sklearn import preprocessing
```

```
In [16]: data["age_scaled"] = preprocessing.scale(data.age.values)
```

**Contents**

- 1 Executive Summary
- 2 Introduction
- 3 Importing Libraries
- 4 Loading Data
- 5 EDA
- ▼ 6 Construct a Linear Regression Model
  - 6.1 Residual Plot
  - 6.2 Check in the Independence
- ▼ 7 Modelling (treating) County as a Fixed E
  - 7.1 Checking the Coefficients
- 8 Build a Mixed Effect Model
- 9 The Nested Random Effects
- 10 Conclusion


In [17]: `data.head(5)`

```
Out[17]:
```

	bounce_time	age	county	location	age_scaled
0	165.548520	16	devon	a	-1.512654
1	167.559314	34	devon	a	-0.722871
2	165.882952	6	devon	a	-1.951423
3	167.685525	19	devon	a	-1.381024
4	169.959681	34	devon	a	-0.722871

In [19]: `sns.distplot(data.age_scaled)`


Out[19]: `<matplotlib.axes._subplots.AxesSubplot at 0x14168e8fb88>`



first we check if the "bounce time" is dependent on the "age", later we will fit a linear regression model

In [21]: `sns.lmplot(x="age", y="bounce_time", data = data)`

Out[21]: `<seaborn.axisgrid.FacetGrid at 0x1416ab0ef88>`




## 6 Construct a Linear Regression Model

```
In [25]: from sklearn.linear_model import LinearRegression
model = LinearRegression(fit_intercept=True)
x= data.age_scaled
y= data.bounce_time
model.fit(x[:, np.newaxis], y)
xfit = np.linspace(-3, 3, 1000)
yfit = model.predict(xfit[:, np.newaxis])
plt.plot(xfit, yfit)
plt.scatter(x, y)
```

Out[25]: `LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)`

```
In [29]: xfit = np.linspace(-3, 3, 1000)
yfit = model.predict(xfit[:, np.newaxis])
plt.plot(xfit, yfit)
plt.scatter(x, y)
```

Out[29]: `<matplotlib.collections.PathCollection at 0x1416bda86c8>`



```
In [35]: print (model.coef_[0])
print (model.intercept_)
```

```
6.279602007970821
201.31646151854164
```

Contents

- 1 Executive Summary
- 2 Introduction
- 3 Importing Libraries
- 4 Loading Data
- 5 EDA
- 6 Construct a Linear Regression Model
  - 6.1 Residual Plot
  - 6.2 Check in the Independence
- 7 Modelling (treating) County as a Fixed Effect
  - 7.1 Checking the Coefficients
- 8 Build a Mixed Effect Model
- 9 The Nested Random Effects
- 10 Conclusion

```
In [44]: y_predict = model.predict (x.values.reshape(-1,1))
RMSE = sqrt((y-y_predict)**2).values.mean()
results = pd.DataFrame()
results["Method"] = ["Linear Regression"]
results["RMSE"] = RMSE
results
```

```
Out[44]:
```

	Method	RMSE
0	Linear Regression	14.928334

6.1 Residual Plot

We Regress y on x and then draw a scatterplot of the residuals

```
In [47]: ax = sns.residplot(x= "age_scaled", y= "bounce_time", data = data, lowess = True)
ax.set(ylabel='Observed - Prediction')
plt.show()
```

Main observation: there are more positive residuals than negative residuals at the highest and lowest predicted value ranges

6.2 Check in the Independence

we compare the bounce times for each county

```
In [48]: sns.catplot(x="county", y="bounce_time", data=data, kind = "swarm")
```

```
Out[48]: <seaborn.axisgrid.FacetGrid at 0x1416bc7fcc8>
```

**Conclusion:** there is substantial grouping so our data is not independent, and thus it is inappropriate to use a linear model for this data.

**Solution:** separate linear regression models for each county but we have to estimate a slope and intercept parameter for each regression and also it effectively reducing our sample sizes for each category.

7 Modelling (treating) County as a Fixed Effect

```
In [62]: counties= data.county.unique()
```

```
In [53]: data_new = pd.concat([data, pd.get_dummies(data.county)], axis = 1)
```

```
In [54]: data_new
```

```
Out[54]:
```

	bounce_time	age	county	location	age_scaled	cheshire	cumbria	devon	dorset	essex	kent	london	norfolk
0	165.548520	16	devon	a	-1.512654	0	0	1	0	0	0	0	0
1	167.559314	34	devon	a	-0.722871	0	0	1	0	0	0	0	0
2	165.882952	6	devon	a	-1.951423	0	0	1	0	0	0	0	0
3	167.685525	19	devon	a	-1.381024	0	0	1	0	0	0	0	0
4	169.959681	34	devon	a	-0.722871	0	0	1	0	0	0	0	0
...	...	...	...	...	...	...	...	...	...	...	...	...	...
475	211.153312	82	essex	c	1.383217	0	0	0	0	1	0	0	0
476	213.577174	59	essex	c	0.374050	0	0	0	0	1	0	0	0
477	207.625105	69	essex	c	0.812818	0	0	0	0	1	0	0	0
478	198.252773	75	essex	c	1.076080	0	0	0	0	1	0	0	0
479	208.055977	66	essex	c	0.681188	0	0	0	0	1	0	0	0

480 rows x 13 columns

Contents

1 Executive Summary

2 Introduction

3 Importing Libraries

4 Loading Data

5 EDA

▼ 6 Construct a Linear Regression Model

6.1 Residual Plot

6.2 Check in the Independence

▼ 7 Modelling (treating) County as a Fixed E

7.1 Checking the Coefficients

8 Build a Mixed Effect Model

9 The Nested Random Effects

10 Conclusion

In [73]:

```
model = LinearRegression(fit_intercept=True)
x = data_new.loc[:,np.concatenate([["age_scaled"],counties])]
y = data.bounce_time
```

In [74]:

```
model.fit(x, y)
```

Out[74]:

```
LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

In [75]:

```
y_predict = model.predict(x)
RMSE = sqrt(((y-y_predict)**2).values.mean())
results.loc[1] = ["Fixed", RMSE]
results
```

Out[75]:

	Method	RMSE
0	Linear Regression	14.928334
1	Fixed	8.563396

7.1 Checking the Coefficients

coefficients of age and the counties

In [76]:

```
pd.DataFrame.from_records(list(zip(np.concatenate([["age_scaled"],counties]), model.coef_)))
```

Out[76]:

	0	1
0	age_scaled	0.048782
1	devon	-21.381957
2	cumbria	9.391460
3	norfolk	9.824419
4	kent	7.938668
5	dorset	-2.079637
6	london	-21.437323
7	cheshire	18.372916
8	essex	-0.628546

**Conclusion:** The residuals are much better than before (more evenly distributed with respect to age). Coefficient for the gradient given to age is substantially smaller, and is likely no longer significant.

**Solution:** we need to control for the variation between the different counties, we have to treat our counties as random effects. So in this new model we treat age (what we are interested in) as a fixed effect, and county and location as a random effect.

8 Build a Mixed Effect Model

First we look at how the bounce time relates to the scaled ages, while controlling for the impact of counties by allowing for a random intercept for each country (each county has its own random intercept, but that the slopes are still the same with respect to age)

In [78]:

```
import statsmodels.api as sm
import statsmodels.formula.api as smf
md = smf.mixedlm("bounce_time ~ age_scaled", data, groups=data["county"])
mdf = md.fit()
print(mdf.summary())
```

Mixed Linear Model Regression Results

=====

Model: MixedLM Dependent Variable: bounce\_time

No. Observations: 480 Method: REML

No. Groups: 8 Scale: 74.7350

Min. group size: 60 Log-Likelihood: -1733.0397

Max. group size: 60 Converged: Yes

Mean group size: 60.0

-----

Coef. Std.Err. z P>|z| [0.025 0.975]

-----

Intercept 201.316 5.175 38.902 0.000 191.174 211.459

age\_scaled 0.136 0.612 0.221 0.825 -1.065 1.336

Group Var 212.999 13.382

=====

In [79]:

```
y_predict = mdf.fittedvalues
RMSE = sqrt(((y-y_predict)**2).values.mean())
results.loc[2] = ["Mixed", RMSE]
results
```

Out[79]:

	Method	RMSE
0	Linear Regression	14.928334
1	Fixed	8.563396
2	Mixed	8.563948

**Conclusion:** The residuals plot looks alomst identical to the previous one where we were treating the county as a fixed effect.

**Solution:** To ensure that each county has its own random slope we need to include this in our random effects formla

Contents

- 1 Executive Summary
- 2 Introduction
- 3 Importing Libraries
- 4 Loading Data
- 5 EDA
- 6 Construct a Linear Regression Model
  - 6.1 Residual Plot
  - 6.2 Check in the Independence
- 7 Modelling (treating) County as a Fixed Effect
  - 7.1 Checking the Coefficients
- 8 Build a Mixed Effect Model
- 9 The Nested Random Effects
- 10 Conclusion

```
In [87]: md = smf.mixedlm("bounce_time ~ age_scaled", data, groups=data["county"], re_formula="~age_scaled")
mdf = md.fit()
print(mdf.summary())
```

Mixed Linear Model Regression Results						
=====						
Model:	MixedLM	Dependent Variable:		bounce_time		
No. Observations:	480	Method:		REML		
No. Groups:	8	Scale:		72.8722		
Min. group size:	60	Log-Likelihood:		-1733.3946		
Max. group size:	60	Converged:		Yes		
Mean group size:	60.0					
-----						
	Coef.	Std.Err.	z	P> z	[0.025	0.975]
-----						
Intercept	202.140	8.356	24.190	0.000	185.762	218.518
age_scaled	0.161	1.196	0.134	0.893	-2.184	2.505
Group Var	558.143					
Group x age_scaled Cov	-51.614					
age_scaled Var	8.621					
=====						

Conclusion: The mixed model with the random slopes is now performing much better, with the residuals much better ditributed. Crucially though, we can see that age does not impact the bounce rate.

9 The Nested Random Effects

The random effects are sometimes nested. For example, there is nothing important about the locations a, b, and c that link location a in one county (London) with that in others (Essex). Therefore explicitly nest these two features.

```
In [81]: data["location_county"] = data["location"] + "_" + data["county"]
```

```
In [82]: data.head()
```

```
Out[82]:
```

	bounce_time	age	county	location	age_scaled	location_county
0	165.548520	16	devon	a	-1.512654	a_devon
1	167.559314	34	devon	a	-0.722871	a_devon
2	165.882952	6	devon	a	-1.951423	a_devon
3	167.685525	19	devon	a	-1.381024	a_devon
4	169.959681	34	devon	a	-0.722871	a_devon

```
In [88]: md = smf.mixedlm("bounce_time ~ age_scaled", data, groups=data["location_county"], re_formula="~age_scaled")
mdf = md.fit()
print(mdf.summary())
```

Mixed Linear Model Regression Results						
=====						
Model:	MixedLM	Dependent Variable:		bounce_time		
No. Observations:	480	Method:		REML		
No. Groups:	24	Scale:		23.7942		
Min. group size:	20	Log-likelihood:		-1504.9078		
Max. group size:	20	Converged:		No		
Mean group size:	20.0					
-----						
	Coef.	Std.Err.	z	P> z	[0.025	0.975]
-----						
Intercept	201.491	3.448	58.441	0.000	194.734	208.249
age_scaled	0.151	0.393	0.385	0.700	-0.618	0.920
Group Var	282.769	21.275				
Group x age_scaled Cov	-8.285	2.478				
age_scaled Var	0.386	0.494				
=====						

```
In [85]: y_predict = mdf.fittedvalues
RMSE = sqrt(((y-y_predict)**2).mean())
results.loc[3] = ["Nested_Mixed", RMSE]
results
```

```
Out[85]:
```

	Method	RMSE
0	Linear Regression	14.928334
1	Fixed	8.563396
2	Mixed	8.563948
3	Nested_Mixed	4.764192

10 Conclusion

We showed that by paying attention to a number of assumptions about the data (homoscedastic and Independence), the interpretaions of the obtained results out of a simple linear regression model could be different. That is to say such investigations can guide us towards building fixed or mixed effect models (sometimes called "multilevel models" or "hierarchical models").