

Contents 🔄 ⚙️

- 1 Executive Summary
- 2 Introduction
- 3 Create a Dataproc Cluster
- 4 Using Spark-bigquery-connector
- 5 Setup a Session
- 6 Get the Data as a Dataframe
- 7 Load Libraries
- 8 Read the Data
- 9 EDA
- 10 Create a map of Italy
- 11 Create a heatmap
- 12 Conclusion

1 Executive Summary

In this notebook I aim to create a data processing pipeline using Apache Spark with Dataproc on Google Cloud Platform. It is a common use case in data science to read data from one storage location (in our case Public Data available on Google cloud), perform transformations on it and write it into another storage location. I will explain how to connect pyspark to BigQuery, get the big data and transform it to a dataframe save it as a CSV file and perform a nice visualization using **"folium"** and **"geopy"** Libraries.

2 Introduction

I plan to build a heatmap for the number of identified Covid19 cases in Italy during '2020-02-24' ~'2020-09-25'. This data is publicly available in google cloud but we aim to have access to that through pyspark. I aim to visualize my findings a heatmap on the map of Italy.

3 Create a Dataproc Cluster

First I create a Dataproc Cluster and I install and run a Jupyter notebook on it. I also enabled the required APIs such as BigQuery API, BigQuery Storage API, etc <https://cloud.google.com/dataproc/docs/tutorials/jupyter-notebook> (<https://cloud.google.com/dataproc/docs/tutorials/jupyter-notebook>) <https://medium.com/google-cloud/apache-spark-and-jupyter-notebooks-made-easy-with-dataproc-component-gateway-fa91d48d6a5a> (<https://medium.com/google-cloud/apache-spark-and-jupyter-notebooks-made-easy-with-dataproc-component-gateway-fa91d48d6a5a>)

4 Using Spark-bigquery-connector

Python 3 kernel within my Dataproc Cluster allow me to configure the SparkSession in my notebook . I only need to include the spark-bigquery-connector required to use the BigQuery Storage API.

5 Setup a Session

```
In [ ]: from pyspark.sql import SparkSession
spark = SparkSession.builder \
    .appName('Myapp') \
    .config('spark.jars', 'gs://spark-lib/bigquery/spark-bigquery-latest_2.12.jar') \
    .getOrCreate()
```

6 Get the Data as a Dataframe

```
In [ ]: table = "bigquery-public-data.covid19_italy.data_by_province"
df = spark.read \
    .format("bigquery") \
    .option("table", table) \
    .load()
df.printSchema()
```

```
In [ ]: # Creating two dataframes
df1=df \
    .select('province_code', 'latitude', 'longitude')
df_agg = df \
    .select('province_code', 'province_name', 'confirmed_cases', 'latitude', 'longitude', 'date') \
    .where("date > '2020-02-24'") \
    .groupBy('province_code') \
    .sum('confirmed_cases')
```

```
In [ ]: # Removing the duplicate records out of the dataframes and joining them
df1 = df1.drop_duplicates(subset=['province_code'])
inner_join = df1.join(df_agg, df1.province_code == df_agg.province_code)
inner_join = inner_join.drop_duplicates(subset=['latitude', 'longitude'])
```

```
In [ ]: #Saving the result
inner_join.toPandas().to_csv('export_dataframe.csv', index = False, header=True)
```

At this point our dataframe is saved on the cloud and we can download it easily to our local system

7 Load Libraries

```
In [22]: import pandas as pd
import numpy as np
import folium
import requests
from geopy.geocoders import Nominatim
from pandas.io.json import json_normalize
```

Folium is a library for creating interactive maps.

8 Read the Data

Contents

1 Executive Summary

2 Introduction

3 Create a Dataproc Cluster

4 Using Spark-bigquery-connector

5 Setup a Session

6 Get the Data as a Dataframe

7 Load Libraries

8 Read the Data

9 EDA

10 Create a map of Italy

11 Create a heatmap

12 Conclusion

```
In [79]: df_train = pd.read_csv('Downloads/export_dataframe.csv')
df_train

Out[79]:
```

	province_code	province_name	latitude	longitude	province_code.1	sum(confirmed_cases)
0	96	Biella	45.566511	8.054082	96	182372
1	23	Verona	45.438390	10.993527	23	912293
2	87	Catania	37.502878	15.087047	87	174827
3	74	Brindisi	40.638485	17.946016	74	113081
4	99	Rimini	44.060901	12.565630	99	400671
...
103	61	Caserta	41.074659	14.332405	61	114684
104	84	Agrigento	37.309711	13.584575	84	30160
105	28	Padova	45.406930	11.876087	28	756088
106	81	Trapani	38.018501	12.513657	81	30183
107	71	Foggia	41.462269	15.543051	71	211754

108 rows × 6 columns

```
In [80]: df_train.isnull().sum()

Out[80]: province_code      0
province_name      0
latitude          1
longitude          1
province_code.1    0
sum(confirmed_cases) 0
dtype: int64

In [81]: df_train=df_train.dropna()

In [82]: df_train

Out[82]:
```

	province_code	province_name	latitude	longitude	province_code.1	sum(confirmed_cases)
0	96	Biella	45.566511	8.054082	96	182372
1	23	Verona	45.438390	10.993527	23	912293
2	87	Catania	37.502878	15.087047	87	174827
3	74	Brindisi	40.638485	17.946016	74	113081
4	99	Rimini	44.060901	12.565630	99	400671
...
103	61	Caserta	41.074659	14.332405	61	114684
104	84	Agrigento	37.309711	13.584575	84	30160
105	28	Padova	45.406930	11.876087	28	756088
106	81	Trapani	38.018501	12.513657	81	30183
107	71	Foggia	41.462269	15.543051	71	211754

107 rows × 6 columns

```
In [83]: df_train= df_train.drop("province_code.1", axis= 1)
```

It seems through we have access to 107 province in Italy only.

9 EDA

```
In [86]: # Let us Look at the top 20 in the List
df_train.sort_values('sum(confirmed_cases)', ascending=False).head(20)

Out[86]:
```

	province_code	province_name	latitude	longitude	sum(confirmed_cases)
94	15	Milano	45.466794	9.190347	4169004
75	17	Brescia	45.539931	10.219103	2742099
89	1	Torino	45.073274	7.680687	2657644
41	16	Bergamo	45.694414	9.668425	2553422
72	19	Cremona	45.133367	10.024209	1193419
77	58	Roma	41.892770	12.483667	1126933
52	108	Monza e della Brianza	45.584390	9.273582	994770
79	10	Genova	44.411493	8.932699	951271
51	18	Pavia	45.185093	9.160157	942463
68	37	Bologna	44.494367	11.341721	928093
1	23	Verona	45.438390	10.993527	912293
73	35	Reggio nell'Emilia	44.697353	10.630080	890031
95	22	Trento	46.068935	11.121231	828280
28	33	Piacenza	45.051935	9.692633	807185
105	28	Padova	45.406930	11.876087	756088
8	36	Modena	44.646000	10.926155	719917
39	6	Alessandria	44.912974	8.615401	686924
14	13	Como	45.809991	9.085160	679186
27	34	Parma	44.801074	10.328350	653653
61	98	Lodi	45.314407	9.503721	647071


10 Create a map of Italy

```
In [75]: def generateBaseMap(default_location=[41.90322, 12.49564], default_zoom_start=5):
base_map = folium.Map(location=default_location, control_scale=True, zoom_start=default_zoom_start)
return base_map
```

Contents

- 1 Executive Summary
- 2 Introduction
- 3 Create a Dataproc Cluster
- 4 Using Spark-bigquery-connector
- 5 Setup a Session
- 6 Get the Data as a Dataframe
- 7 Load Libraries
- 8 Read the Data
- 9 EDA
- 10 Create a map of Italy
- 11 Create a heatmap
- 12 Conclusion

In [76]: baseMap = generateBaseMap()
baseMap



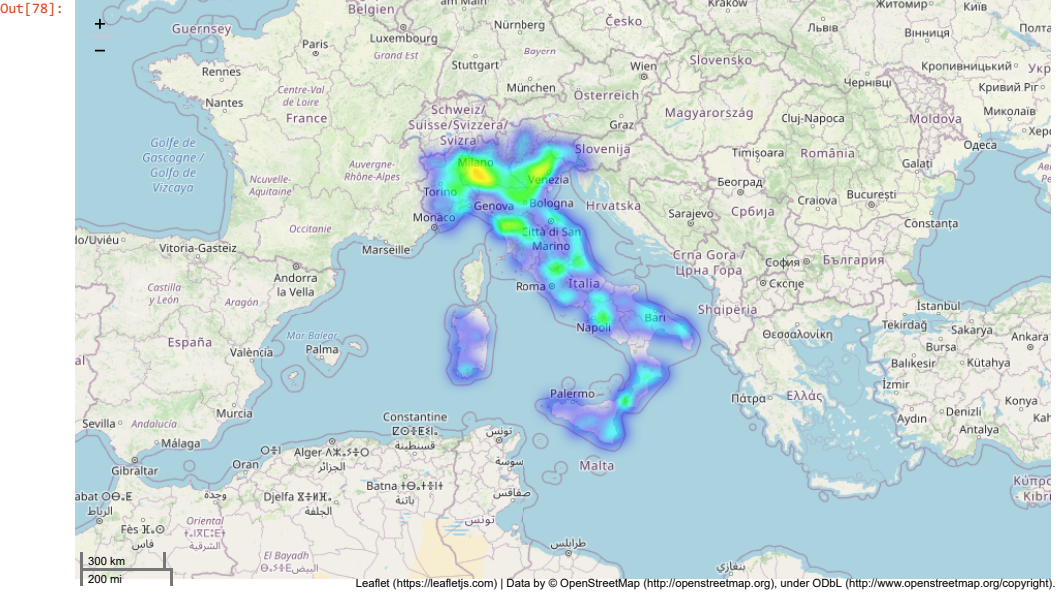
11 Create a heatmap

In [69]: from folium.plugins import HeatMap

In [77]: base_map = generateBaseMap()
HeatMap(data=df_train[["latitude", "longitude", "sum(confirmed_cases)"]].values.tolist(),
radius=8, max_zoom=5).add_to(base_map)

Out[77]: <folium.plugins.heat_map.HeatMap at 0x1a44df3cb88>

In [78]: base_map



12 Conclusion

We did not have access to the region information, but our visualization for sure helps us to see the number of positive cases are outnumbered in the north part of Italy.